



# Research Review of Distributed Multi-Agent Planning

*By Mohamed Bakr Ali*

---

## Introduction:

Having greatly enjoyed working with distributed platforms (Spark, Hadoop, etc.) in the past, for this research review, we chose to focus on the challenges of **distributed multi-agent planning (DMAP)** techniques. To guide us through this quick exploration, we relied in large parts on the work of B.K. Durkota as presented in [1].

## DMAP Challenges:

The challenges of DMAP are many. If a particular agent is unable to solve a problem on its own and has its own actions, **how should its actions and the actions of the other agents be interleaved?** Coordination implies communication. In order not to incur a huge network overhead and slowing down the agents too much, **what is the best paradigm to use for coordination?**

## DMAP Approaches:

- Distributed Constraint Satisfaction Problem solving for coordination of the agents and individual planning using local search [2],
- Multi-agent adaptation of A\* with local heuristics [3],
- Distribution of the GraphPlan approach based on merging of planning graphs [4].

In [2], Nissim et al use the **DisCSP+Planning** algorithm to solve the multi-agent planning problem by separating the **public** and **private** aspects of the problems. In this paradigm, the public aspects (coordination between agents) is dealt with by the **coordination component** using a CSP searching for a sequence of **interaction points** between the agents and enforcing consistency requirements between actions. The local, internal aspect is dealt with by the **individual planning component** using a planner that handles the other types of constraints and encodes the local parts of the plan.

In [3], the authors based their work on the A\* algorithm described in class. Similarly to regular A\*, the **Multi-agent Distributed A\* (MA-A\*)** algorithm maintains open lists of unvisited states and closed lists of already visited states for all agents. Individual agents use local, potentially different, heuristics to decide which unvisited state should

---

expand next. As in [2], actions must first be separated into public and individual actions. Agents send messages to each other to distribute the search at interaction points where the other agents can follow.

In [4], the authors use a distributed version of the planning graph data structure. They use the **Distributed Planning through Graph Merging (DPGM)** algorithm which first performs **global goal decomposition** (where each agent creates an individual goal). Then, it alternates **expansion** (where each agent builds new layers in their planning graphs) and **planning graph merging** (where agents share their actions until each of them reaches their individual goals). In the **individual plan extraction** phase, each agent extracts plans from its planning graph. The process ends with a **coordination** phase that yields a coordinated individual solution plan.

## Results:

In [1], experimental results show DPGM to be efficient in domains which are not tightly coupled ("combinatorically easy"). The DisCSP+Planning algorithm is shown to be efficient in problems which are combinatorically hard from the perspective of individual planning. Finally, the implementation of MA-A\* with set additive heuristics was shown to be most effective in highly coupled domains.

## References:

- [1] *Comparison of Deterministic Distributed and Multi-Agent Planning Techniques* (2013), by B.K. Durkota
- [2] *A general, fully distributed multi-agent planning algorithm* (2010), by R. Nissim et al
- [3] *Multi-Agent A\* for Parallel and Distributed Systems* (2012), by R. Nissim et al
- [4] *Distributed planning through graph merging* (2010), by D. Pellier